

Data Processing on Modern Hardware

Jana Giceva

Lecture 8: Multicore CPUs


NUMA, chiplet-based processors



Non-uniform memory access (NUMA)


Distributed Shared Memory

- Single processor scalability (with shared memory) has limitations
- Idea: distribute memory



NUMA hardware

- Almost all mid-range and enterprise servers today are multi-socket



- Each server typically contains between 2-8 sockets
- Each socket contains:
 - between 4 and 52 cores (up to 96 cores on AMD Zen4)
 - a few memory DIMM modules attached through memory channels
- An interconnect network among the sockets allows each core to access non-local memory


Effects of NUMA Hardware

■ Example multi-socket server:

- four 8-core Nehalem-EX processors, fully connected with 4 bi-directional 3.2 GHz QuickPath Interconnect (QPI)

■ Measure performance for reading local vs. remote memory

- **Flow 1: read locally** (from socket 0)
 - max. aggregate bandwidth (12 threads) is 24.7 GB/s
 - latency is 340 CPU cycles (~ 150ns)
- **Flow 2: read remote** (over **1 QPI link**, from socket 3)
 - max. aggregate bandwidth is 10.9 GB/s
 - latency is 420 CPU cycles (~185ns)
- **Flow 3: read remote** (over **2 QPI links**, from socket 1)
 - max. aggregate bandwidth is 10.9 GB/s
 - latency is 520 CPU cycles (~230ns)
- **Flow 4: read remote** (over **2 QPI links**) **with cross traffic**
 - max. aggregate bandwidth is 5.3 GB/s
 - latency is 530 CPU cycles (~235ns)



What does that mean for data processing?

- **Designing algorithms and data structures**

- Need to differentiate between local and remote memory
- Local memory is faster and has higher bandwidth

- **Concurrency**

- Synchronization within a socket / NUMA node is significantly faster
- Concurrent data structures need to scale across NUMA nodes

- **NUMA effects in systems and databases**

System's support for NUMA

- Modern **operating systems** are aware of NUMA architectures.
 - Linux partitions memory into NUMA zones, one for each socket.
 - For each NUMA zone, the kernel maintains separate management data structures.
- By **default** the Linux kernel **allocates memory** on the local NUMA node
 - The socket of the core on which the current thread is scheduled on.
- **Unless explicitly bound** to a specified socket (NUMA zone) through the **mbind()** system call

■ Watch out for default OS

- **First touch** allocation **policy** (static and in place **until kernel version 2.6**)
- **Today** there are two options:
 - **Transparent NUMA awareness:**
 - Allocate locally
 - Migrate thread or data to achieve good NUMA performance and balancing
 - **Explicit memory allocation policy**
 - Allocate memory (and do not migrate) based on the selected policy

■ NUMA memory policy

- System default policy
 - local (general), interleaved (during boot-up)
- Task/Process policy – controls all page allocations made by or on behalf of the task
- VMA policy – to a range of a task's virtual address space

■ Allocation modes

- **Local** – memory from the local NUMA node
- **Bind** – memory from the set of nodes specified by the policy
- **Preferred** – memory from the set of nodes specified by the policy, if available
- **Interleaved** – memory interleaved across all the NUMA nodes in the set provided by the policy

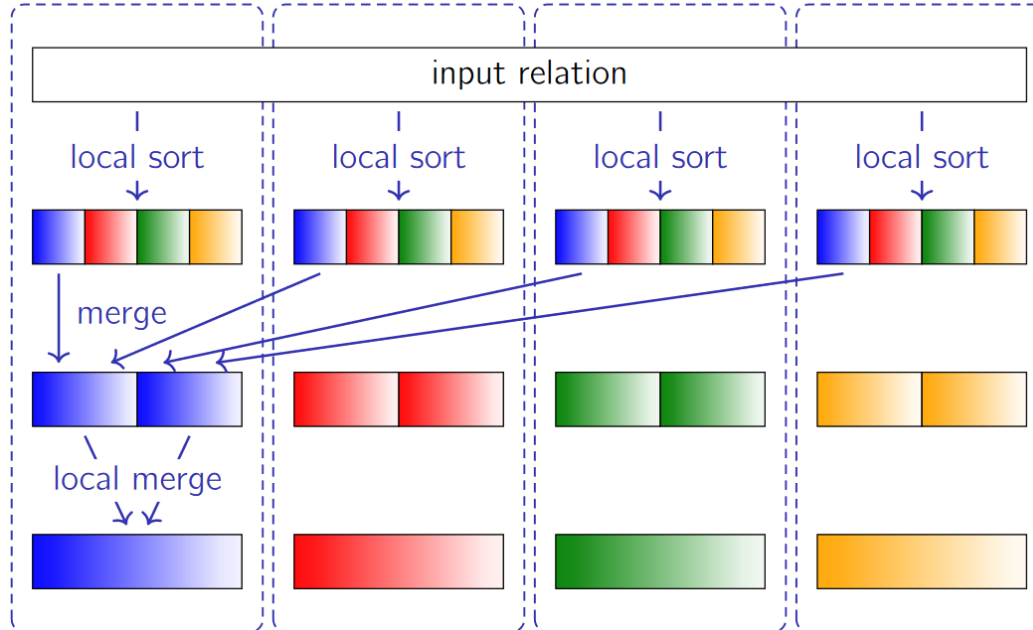
■ Invoked from the process / thread itself or via the numactl library

■ Memory policy APIs

- `long set_mempolicy(int mode, const unsigned long *nmask, unsigned long maxnode)`
- `long get_mempolicy(int *mode, const unsigned long *nmask, unsigned long maxnode, void *addr, int flags);`
- `long mbind(void *start, unsigned long len, int mode, const unsigned long *nmask, unsigned long maxnode, unsigned flags)`

Where does it matter?

■ Example 1: Sorting and NUMA




Step 1: Operate on local data as much as possible.

Step 2: Exchange data to gather local partitions on a local place.

Step 3: Merge the results locally.

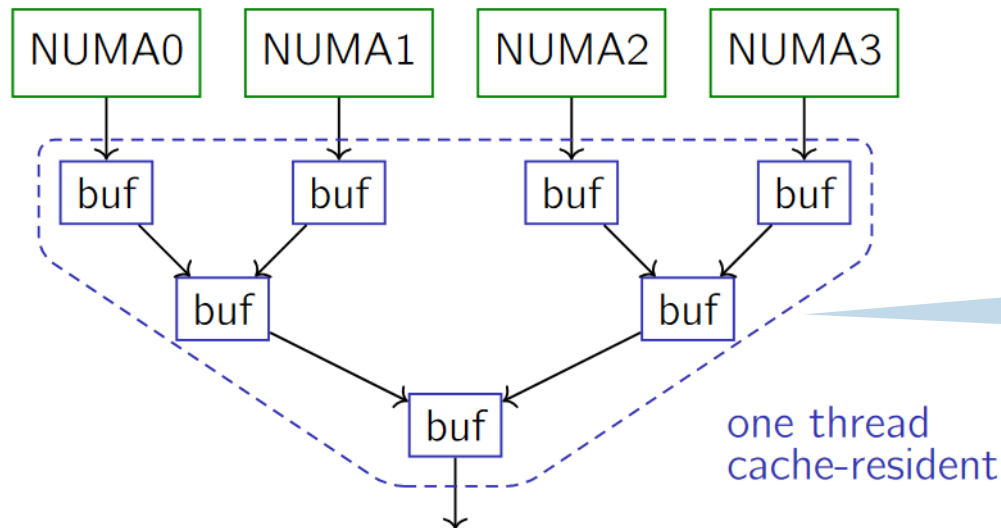
Sorting and NUMA



Even though NUMA-aware, the algorithm does not scale well.

Step 2 is very memory-bandwidth intensive and saturates the system's resources.


Multi-way merging as an alternative



Step 1: Operate on local data as much as possible. (similar to before)

Step 2: Gather data in a cache- and NUMA-conscious way. Recall the multi-way sort when we did SIMD?

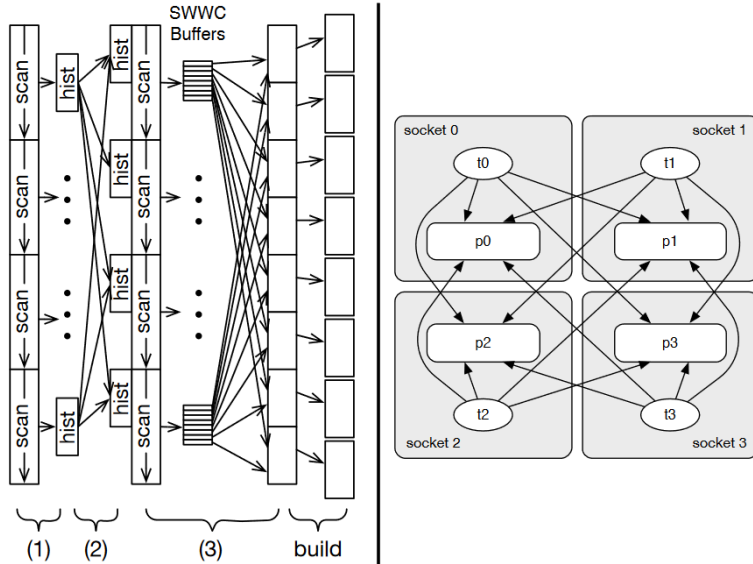
Sorting and NUMA



Performance speed-up much better than before because of NUMA-awareness.


Radix-join and NUMA

v1: radix-join PRO



- Step (3) introduces a lot of small random writes to remote memory (different NUMA regions).
- Can we somehow avoid it?

v2: radix-join CPRL



- Skip step (2) – computing the global histogram
- Proceed with a node-local step (3)
- Trade-off small remote-writes for large remote-reads



Radix-join and NUMA

■ How about the join phase?


- All radix-join algorithms build co-partitions that should be joined independently.
- They are scheduled using a LIFO task-queue, to be processed by different threads.
- With a naive schedule, all the first threads taking a task from the queue will have to read their input data from the same NUMA region.
- For $\frac{3}{4}$ of the threads, this is a remote node.

■ NUMA-aware scheduling

- Can improve this by reordering the join tasks in an interleaved fashion or by using separate task queues for each NUMA region.



Radix-join and NUMA



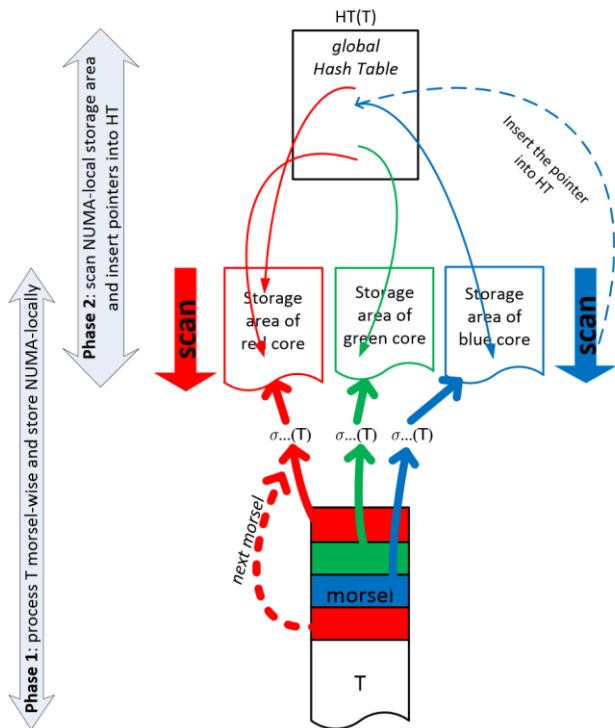



Figure 3: NUMA-aware processing of the build-phase


- Relation T is interleaved “morsel-wise” across the NUMA nodes
- The scheduler assigns a morsel located on the same NUMA node where the thread is executed.
- In the first phase the filtered tuples are inserted into NUMA-local storage areas, i.e., for each core there is a separate storage area in order to avoid synchronization.
- The global HT is probed by threads located on various sockets of a NUMA system.
 - To avoid contention, it is interleaved across all sockets.

Engine-wide NUMA awareness



Impact of NUMA on DB synchronization

- Performance implications for synchronization and concurrency



- **Goal:** check the impact of NUMA latencies on OLTP transactions and the overall throughput.
- **OLTP workload:** TPC-C payment transaction
- **Machine:** 4 CPUs with 6 cores each
- **Test:** Run the database with 4 worker threads, either using the default OS scheduling or pinning them to different cores.
- **Insights:**
 - DB threads collocated on the same NUMA node exhibit much better performance than alternatives.
 - Communication over the interconnect is expensive.
 - OS-scheduling can be unpredictable.

- **Synchronization within the processor is cheaper than to synchronization over the interconnect**
 - due to **latency** concerns, but also for **increased memory traffic**.
- Two main approaches to make locks NUMA-aware locks (and concurrent data structures):
 - **Cohort locks** (hierarchical locks) [1]
 - **Combining** + remote core execution (select a leader, etc.) [2]
- Recent **black box approach** allows any **linear data structure** to be made NUMA-aware [3]
- Parking lock (e.g., optimized futex from last week) can be made a scalable and NUMA-aware blocking synchronization primitive: CST [4]

[1] Chabbi et al. *High Performance Locks for Multi-Level NUMA Systems*. PPOPP 2015

[2] Lozi et al. *Fast and Portable Locking for Multicore Architectures*. ACM Trans. Computing Systems 2016

[3] Calciu et al. *Black-box Concurrent Data Structures for NUMA Architectures*. ASPLOS 2017

[4] Kashyap et al. *Scalable NUMA-aware Blocking Synchronization Primitives*. USENIX ATC 2017

<https://taesoo.kim/pubs/2017/kashyap:cst-slides.pdf>

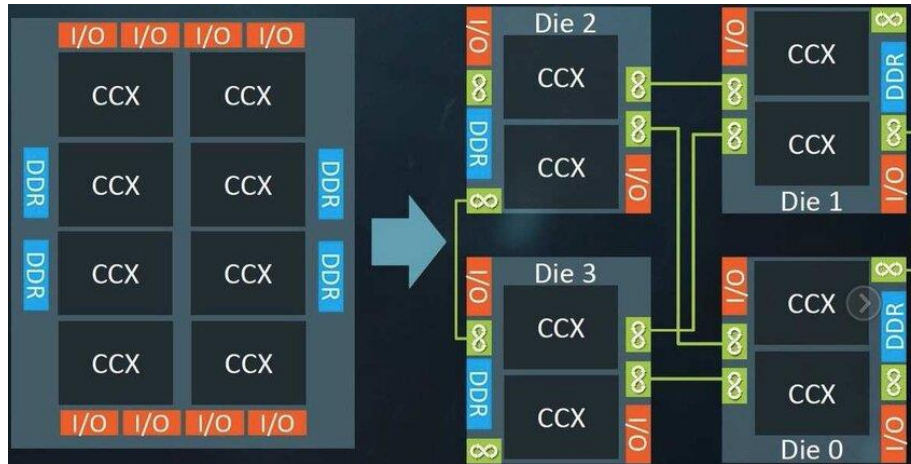
Chiplet-based processors

- A *chiplet* is a small, specialized die (e.g., CPU core cluster, GPU, I/O) designed to work as part of a larger system.
 - Lego-like approach that let's designers mix and match dies (even from different vendors) to build more complex SoCs cost-effectively and chiplet-independent innovation.
- Technology tailoring: each chiplet can use the optimal process node for its function
 - Cutting-edge 3-5nm for CPU logic
 - Mature 6-28nm for I/O
 - Higher yield → chance of defect killing a die drops with its area
- Die-to-die links:
 - High-speed serial links connect the chiplets. Often packetized (unlike on-chip buses) and are thus designed for bandwidth.
 - Examples: AMD's Infinity Fabric, Intel's Foveros interconnect, NVLink

Chiptlets vs. Monolithic Processors

- A **monolithic CPU/GPU** places **all** logic on one large die
- A **chiptlet** splits functions into **separate** dies.

https://www.researchgate.net/figure/Monolithic-Dieleft-vs-Chiptletsright_fig4_376192917



- Pros chiptlets:
 - Higher yield, lower per-chip wafer cost
 - Better scalability to many cores/dies
 - E.g., AMD scaled EPYC from 8-core to 192-cores by adding chiptlets (otherwise impractical)
- Cons chiptlets:
 - Inter-chiplet communication is slower


■ Performance/power trade-off

- Monolithic chips achieve low latency for tight coupling
- Chiptlets introduce latency/bandwidth penalty across the interconnect

Examples

■ AMD EPYC (Zen Architecture)


- Use multi-chip modules chiplet architectures
- Each Core Complex Die (CCD) is an 8-core Zen chiplet
- A central memory I/O die (IOD) has the memory controllers, PCIe, etc.
- CCDs are placed around the IOD and connected with the Infinity Fabric.



Examples

■ Intel Sapphire Rapids XCC (tile design)


- Four tiles per package
- Each tile with up to 15 cores, and a single UPI link
- Two channels for DDR5
- Accelerators:
 - Data Streaming Accelerator
 - Quick Assist Technology
 - Dynamic Load Balancing
 - In-Memory Analytics Accelerator
- Compute Express Link (CXL 1.1)



Other examples

■ ARM Graviton 3

- 3 sockets connected and managed with one Nitro card
- Each socket CPU follows a chiplet design with 7 silicon dies
- All 64 cores placed on a single chiplet
- Remaining 6 chiplets handle IO, DDR5/PCIe 5.0



■ Apple M1 Ultra

- Glues two M1 Max dies side-by-side
- Each M1 includes CPU cores, GPU cores, an NPU and a unified physical pool of DDR memory
- The result acts like one 20-core CPU, 64-core GPU and 128 GB unified RAM.



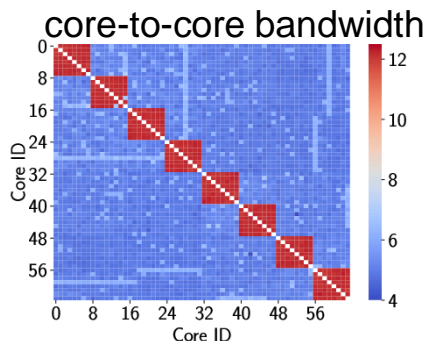


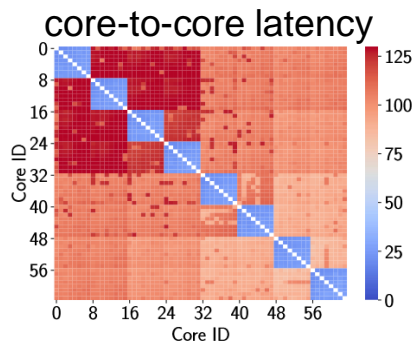
■ NVIDIA Grace Hopper (Het. Superchip)

- CPU+GPU Superchip
 - Grace CPU die (ARM Neoverse cores, up to 512 GB LPDDR5X 546 Gb/s), Hopper GPU die (96 GB HBM, 3000 GB/s)
 - Connected by NVLink C2C interconnect fully cache-coherent (900 GB/s).

System Design effect on properties

■ NUMA-like effects:

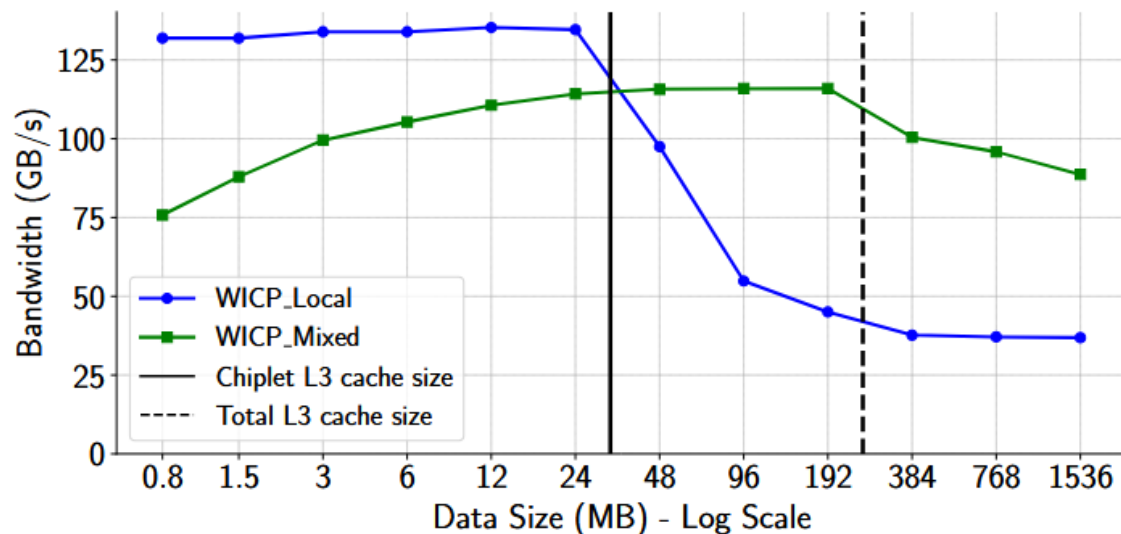
- multi-chiplet designs often behave like NUMA systems even on a single socket
- Each chiplet is a “node” with local memory (except ARM’s Graviton)
- Accessing another chiplet’s memory adds latency and reduces bandwidth



- Q: on the right we see the behavior of three architectures: AMD, ARM and Intel. Which is which?

The effect of partitioned LLC cache

- The visible LLC does not expose that it is partitioned among the chiplets
- Yet, it can have a significant effect on our program's performance



As long as the working set size fits the L3 chiplet portion, keep the threads working together on that chiplet.

Beware of the cliff – locality brings performance but overoptimizing may hurt us once the working set exceeds the threshold. Then spreading is better.


Is it the same across architectures?

- Our experience shows that there is a performance difference between architectures, but the principles apply.


AMD has more chiplets and achieves higher aggregate bandwidth.

For Intel, the difference is not as obvious – expect less speed-up, but also the aggregate bandwidth is considerably lower.


Chiplet-awareness has no effect on monolithic systems.




(a) AMD EPYC Milan



(c) ARM Graviton 3



(b) Intel Sapphire Rapids

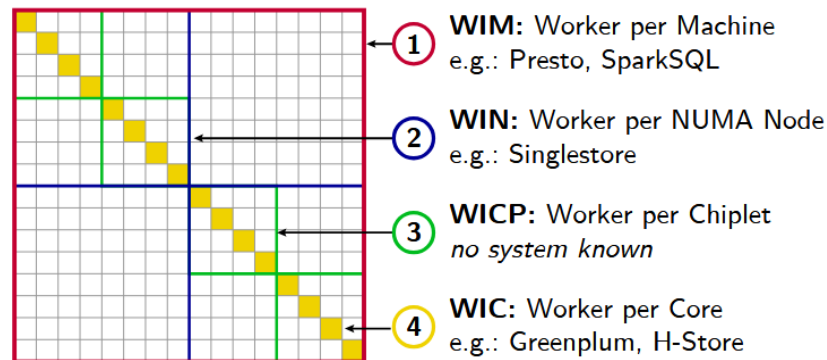


(d) 3rd Gen. Intel Xeon Gold


Chiplet-aware scheduling

■ OLAP on modern Chiplet-based Processors

- Default: one worker per machine
- Optimized: one worker per NUMA
- Some can't leverage parallel execution:
 - One worker per core
- Not explored: one worker per chiplet




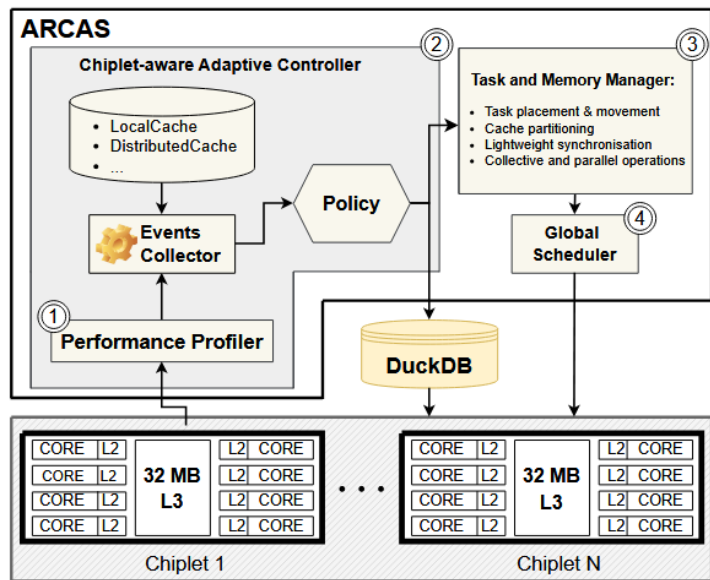
■ Results are quite surprising



Chiplet-aware scheduling

■ Q: Does it generalize beyond analytics?

■ Yes, and we can even make it adaptive.



- Light-weight threads – co-routines
- Task-based scheduling
- Light-weight monitoring on local/remote chiplet accesses
- Runtime adaptive migration of tasks to workers, depending on the working set size

Does it affect how we design algorithms?

■ Let's look at a **sorting**


– **Cache-conscious sort**

- In-register handles runs that fit within the SIMD
- In-cache includes runs that are confined within the CPU's L3 cache
- Out-of-cache is for runs that exceed the cache capacity.

– **Chiplets now partition the L3 cache and thus require rethinking the in-cache sorting:**

- In-local-chiplet-cache vs.
in-remote-chiplet cache.

– **What about the aggregate bandwidth?**



From NUMA to chiplet-aware sorting

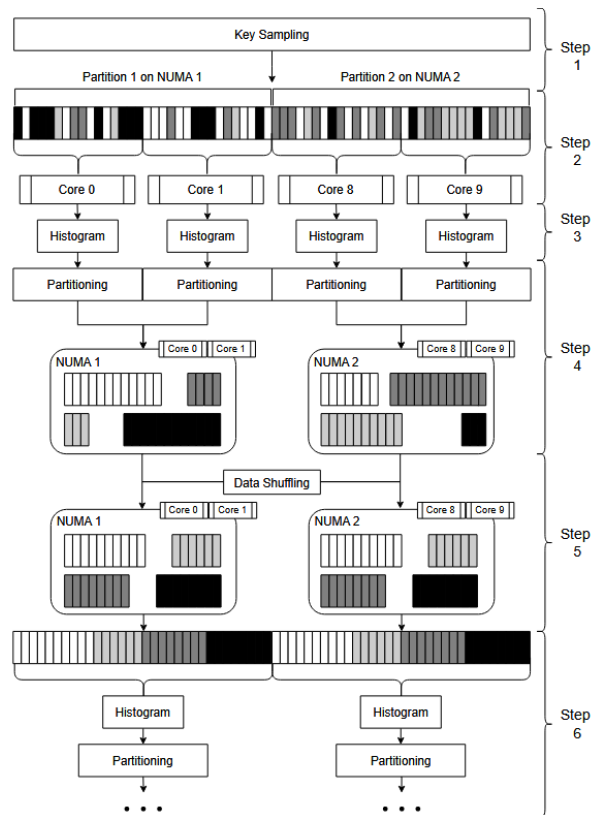


Fig. 5: NUMA-aware LSB Radix-Sort step by step execution.

Step 1: from NUMA-balanced partitions to degree of parallelism

Step 2: bind each part to a NUMA node. Now, bind each thread to a chiplet. Rely on chiplet-aware task-scheduling.

Step 3: unchanged

Step 4-5: avoid the NUMA data shuffling.

Step 6: the process is repeated based on the sorting bits.

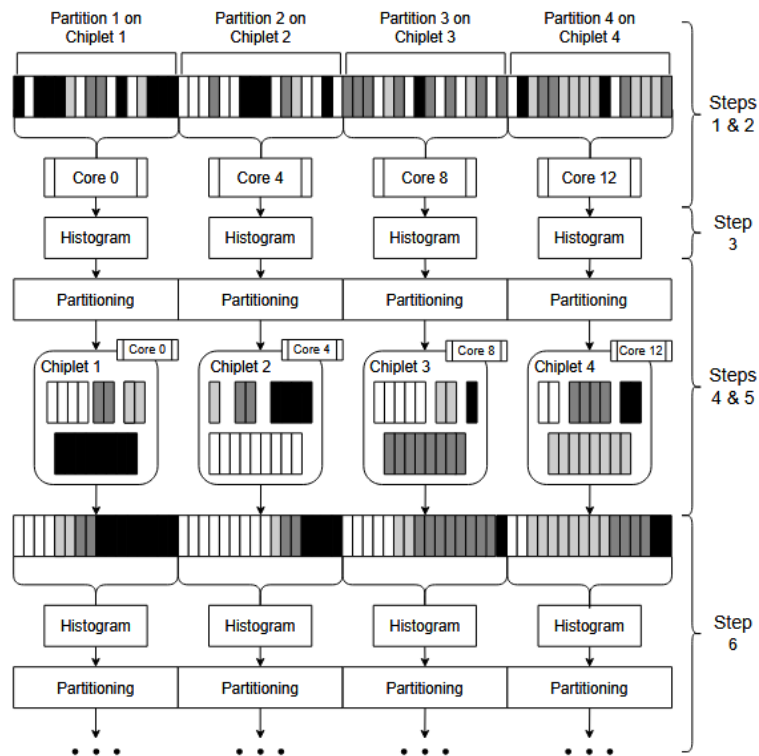




Fig. 6: Chiplet-aware LSB Radix-Sort step by step execution.

Performance impact on Radix-sort



(a) 16 Bit Radix-Sort, 100M Tuples




(a) Radix-Sort


- Key take-away:
 - Chiplet-aware sorting gives a significant boost to NUMA-optimized alternatives
 - One observes approximately 40% performance improvements.
- Similar observations on various dataset sizes, 16-,32-,64-bit keys, and skew distributions.

Why the performance boost?


- Partly by removing the data shuffling phase between the NUMA nodes.
- Partly by adjusting the number of partitioning phases by looking at the data size.
- And finally, by adjusting between chiplet-local vs. mixed on aggregate bandwidth.



(a) NUMA-aware LSB Radix-Sort



(b) Chiplet-aware LSB Radix-Sort



- Various papers cross-referenced in the slides
 - Li et al. “NUMA-aware algorithm: the case of data shuffling” *CIDR 2013*
 - Balkesen et al. “Multi-core, Main-Memory Joins: Sort vs Hash Revisited” *VLDB 2014*
 - Schul et al. “An Experimental Comparison of Thirteen Relational Equi-Joins in Main Memory.” *SIGMOD 2016*
 - Leis et al. “Morse-Driven Parallelism: A NUMA-aware query evaluation framework for the many-core age” *SIGMOD 2014*
 - Porobic et al. “Analyzing the impact of system architecture on the scalability of OLTP engines for high-contention workloads” *VLDB 2017*
 - Lee et al. “MCC-DB: Minimizing Cache Conflicts in Multi-core Processors for Databases” *VLDB 2009*
 - Lo et al. “Heracles: Improving Resource Efficiency at Scale” *ISCA 2015*
 - Makreshanski et al. “BatchDB: Efficient Isolated Execution of Hybrid OLTP and OLAP workloads for Interactive Applications” *SIGMOD’17*
- Lecture: *Data Processing on Modern Hardware* by Prof. Jens Teubner (TU Dortmund, past ETH)
- Book: *What every programmer should know about memory?* by Ulrich Drepper
 - Chapters 5 and 6.5
- Intel Architectures Software Developer Manuals
 - Optimizing Applications for NUMA (<https://software.intel.com/content/dam/develop/external/us/en/documents/3-5-memmgmt-optimizing-applications-for-numa-184398.pdf>)
 - Volume 3b: chapters 17.16 and 17.16 (for Intel RDT)