

# Data Processing on Modern Hardware

## Tutorial 5

Ferdinand Gruber


Michalis Georgoulakis



# Assignment 5 - Synchronization

# Assignment 5 - Synchronization

- Databases are often faced with highly concurrent workloads
- Hardware offers us parallelization opportunity in multiple cores
- Synchronization :'(




# Assignment 5 - Synchronization

- **Goal:** synchronize access on a list-based set
  - sorted
  - no duplicates
- Supported Operations

```
bool contains(T k) { return false; }  
void insert(T k) { }  
void remove(T k) { }
```

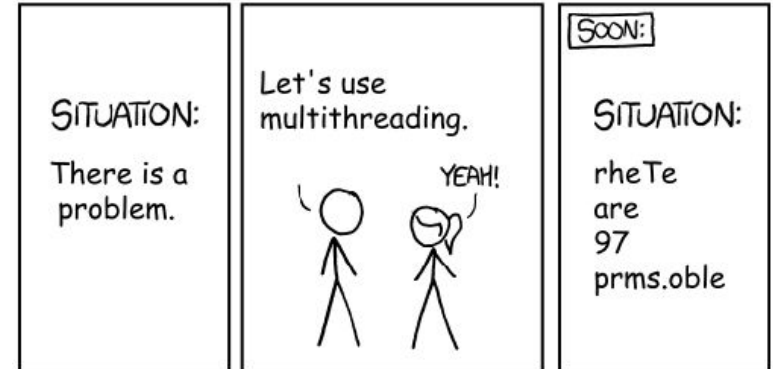
- We give you a baseline implementation without synchronization



# Assignment 5 - Synchronization

You give us implementations of the three operations with the following approaches

- Coarse-Grained Locking
- Coarse-Grained Locking with Read/Write Locks
- Lock Coupling
- Lock Coupling with Read/Write Locks
- Optimistic Locking
- **Bonus:** Optimistic Lock Coupling



# Assignment 5 - Synchronization

## Optimistic Lock Coupling

### traditional

1. lock node A
2. access node A



3. lock node B
4. unlock node A
5. access node B



6. lock node C
7. unlock node B
8. access node C
9. unlock node C



### optimistic

1. read version v3
2. access node A

3. read version v7
4. validate version v3
5. access node B

6. read version v5
7. validate version v7
8. access node C
9. validate version v5

Figure 1: Comparison of a lookup operation in a 3-level tree using traditional lock coupling (left-hand side) vs. **optimistic lock coupling** (right-hand side).

```
struct Entry {  
    T key;  
    std::atomic<Entry *> next;  
    std::atomic<uint64_t> version = 0; // Version counter for the lock coupling  
    M mutex;                          // Mutex for each element in the list  
};
```

# Assignment 5 - Synchronization

## Thread Management tools

### C++

- **TBB Library** => provides abstractions, easier to use (`tbb::spin_mutex`, `tbb::spin_rw_mutex`)
  - Intel library with high-level abstractions for parallelism. It abstracts the complexity of thread management and allows developers to focus on parallel algorithms.
  - **TBB** creates a pool of system threads to execute tasks concurrently (`tbb::task_arena`).

### C

- **Threads (POSIX)** => more control, but also more manual management and synchronization (`pthread_create` / `pthread_join`)

# Assignment 5 - Synchronization

## Mutexes

- Synchronization primitives to protect shared data structures

**Lock (exclusive):** prevents other threads from acquiring the same mutex until it is released.

**Shared Lock:** provides shared access to the resource while preventing exclusive write access by other threads

**Scoped Lock:** allows for automatic locking and unlocking of multiple mutexes in a scoped manner.

### Mutex Flavors

`tbb::spin_mutex` does not scale well, but can be very fast in lightly contended situations



# Assignment 5 - Synchronization

## Analysis

- How expensive is locking for the different approaches?
- How do the approaches perform in regard to #clock-cycles, #instructions and IPC?
- Which approach provides the best performance?

Help: “bench” function in **main.cpp**

## Workload Types

- Read-only workload (contains)
- Mixed workload (insert, update, and contains)

Vary workload type, # of threads, domain name.

# Assignment 5 - Synchronization

Deadline: **11/06**

## Submission Instructions

- First fork the assignment repository.
- Then, in your forked repository, add:
  - Code that implements the assignment
  - A 1-page report answering the assignment questions (**report.pdf**)

Before you get started, you can have a look at:

- TBB Parallel Processing Example
- Optimistic Lock Coupling [Paper](#)
- <https://databasearchitects.blogspot.com/2020/10/c-concurrency-model-on-x86-for-dummies.html>

# **Assignment 3 - Hardware Optimized Hash Joins**

## **Sample Answers**

# Assignment 3

## 1. Hash Join Baseline

```
// Assume relation r is smaller, return number of matched tuples
uint64_t hash_join(relation &r, relation &s) {
    uint64_t matches = 0;
    std::unordered_set<keyType> hashTable;
    // Step 1: build phases
    for (auto &t : r) {
        hashTable.insert(t.key);
    }

    // Step 2: probe phase
    for (auto &t : s) {
        matches += hashTable.count(t.key);
    }
    return matches;
}
```

**Build Phase:** Easily parallelizable

**Probe Phase:** Requires no synchronization

**Data Structures:**

We use `unordered_set`, `unordered_map` to group unique elements.

## Assignment 3

# 2.1 Naive Partitioning

```
partition partition_naive(relation &r, size_t start, size_t end, uint8_t bits, uint8_t shift) {
    SplitHelper split(bits);
    // TODO implement the naive partitioning here

    // Step 1 build histograms -> Prefix sum
    // create a histogram with #entries = #partitions = fanOut
    std::vector<uint64_t> histogram(split.fanOut, 0);
    for (size_t i = start; i < end; i++) {
        auto bucket = (r[i].key >> shift) & split.mask;
        histogram[bucket]++;
    }

    // Step 2 use prefix sum to partition data
    std::vector<uint64_t> startPositions(split.fanOut, 0);
    uint64_t prefixSum = 0;
    for (size_t i = 0; i < split.fanOut; i++) {
        startPositions[i] = prefixSum;
        prefixSum += histogram[i];
    }

    // Step 3 partition
    relation partitionedRelation(prefixSum);
    std::vector<uint64_t> offset(split.fanOut, 0);
    for (size_t i = start; i < end; i++) {
        auto bucket = (r[i].key >> shift) & split.mask;
        auto position = startPositions[bucket] + offset[bucket];
        partitionedRelation[position] = r[i];
        offset[bucket]++;
    }

    return {partitionedRelation, startPositions};
}
```

1. Create histogram equal to the number of partitions.
2. Iterate over histogram to calculate starting position of each partition in the partitioned relation.
3. Bucket number + offset determine the position where each relation should be placed.

## Assignment 3      2.2 Multi-pass Partitioning

```
partition partition_multiPass(relation &r, uint8_t bits1, uint8_t bits2) {
    // Partition 1. stage
    partition p1 = partition_naive(r, bits1);

    // Partition 2. stage
    relation result;
    std::vector<uint64_t> startPositions;
    uint64_t offset = 0;

    for (uint64_t i = 0, limit = p1.s.size(); i < limit; i++) {
        auto start = p1.s[i];
        auto end = (i == limit - 1) ? p1.r.size() : p1.s[i + 1];

        auto p2 = partition_naive(p1.r, start, end, bits2, bits1);

        for (auto &x : p2.r) {
            result.push_back(x);
        }
        for (auto &x : p2.s) {
            startPositions.push_back(offset + x);
        }
        offset += end - start;
    }

    return {result, startPositions};
}
```

**Idea:** Creating too many partitions can easily thrash the TLB cache.

- Each partition requires its own TLB entries.

**Motivation:** Splitting partitioning into two phases can reduce the fan-out of each stage.

## Assignment 3

### 2.3 Software-managed Buffers & Non-Temporal Writes

```
std::pair<partition, std::vector<uint64_t>> partition_softwareManaged(relation &r, uint8_t bits) {
    SplitHelper split(bits);

    // Step 1 build histograms -> Prefix sum
    std::vector<uint64_t> histogram(split.fanOut, 0);
    for (auto &t : r) {
        auto bucket = t.key & split.mask;
        histogram[bucket]++;
    }

    // Step 2 use prefix sum to partition data
    std::vector<uint64_t> startPositions(split.fanOut, 0);
    uint64_t prefixSum = 0;
    for (size_t i = 0; i < split.fanOut; i++) {
        startPositions[i] = prefixSum;
        auto tmp = histogram[i];
        if (tmp % tuplesPerCL != 0) {
            // align the offsets for non-temporal writes
            tmp += tuplesPerCL - (tmp % tuplesPerCL);
        }
        prefixSum += tmp; // update the offset of the next bucket
    }
}
```

Ensure each partition size  
is a **multiple of cache line  
size**.

## Assignment 3

### 2.3 Software-managed Buffers & Non-Temporal Writes

```
// Step 3 partition
relation partitionedRelation(prefixSum);
std::vector<uint64_t> offset(split.fanOut, 0);
// initialize a software managed buffer per partition and its current offsets
std::vector<SoftwareManagedBuffer> softwareManagedBuffers(split.fanOut);
std::vector<uint64_t> bufferOffsets(split.fanOut, 0);

for (size_t i = 0, limit = r.size(); i < limit; i++) {
    auto bucket = r[i].key & split.mask;

    auto &buffer = softwareManagedBuffers[bucket];
    auto &position = bufferOffsets[bucket];

    buffer.tuples[position] = r[i];
    position++;

    // perform non temporal write when the buffer is full
    if (position == tuplesPerCL) {
        auto outPosition = &partitionedRelation[startPositions[bucket] + offset[bucket]];
        auto writePtr = reinterpret_cast<uint8_t *>(outPosition);
        storeNontemp(writePtr, &buffer);
        offset[bucket] += tuplesPerCL;
        position = 0;
    }
}
```

Create and populate software-managed buffers for each partition.

When buffer is full, non-temporal write to memory.



```
// handle non-empty buffers
for (uint64_t i = 0, limit = split.fanOut; i < limit; i++) {
    auto &buffer = softwareManagedBuffers[i];
    auto &position = bufferOffsets[i];

    auto outPosition = startPositions[i] + offset[i];
    for (size_t j = 0; j < position; j++) {
        partitionedRelation[outPosition + j] = buffer.tuples[j];
    }
    offset[i] += position;
}

partition result = {partitionedRelation, startPositions};
return std::make_pair(result, offset);
}
```

Handle non-empty buffers.

# Assignment 3

## 3 Radix Join

```
// Assume relation r is smaller, return number of matched tuples
uint64_t radix_join(relation &r, relation &s, Partitioning p) {
```

```
    uint64_t matches = 0;
```

```
// Step 1: partitioning phase
```

```
    partition pR{};
```

```
    partition pS{};
```

```
    std::vector<uint64_t> pROffsets;
```

```
    std::vector<uint64_t> pSOffsets;
```

```
    switch (p) {
```

```
    case Partitioning::naive:
```

```
        pR = partition_naive(r, 8);
```

```
        pS = partition_naive(s, 8);
```

```
        break;
```

```
    case Partitioning::multiPass:
```

```
        pR = partition_multiPass(r, 2, 2);
```

```
        pS = partition_multiPass(s, 2, 2);
```

```
        break;
```

```
    case Partitioning::softwareManaged:
```

```
        std::tie(pR, pROffsets) = partition_softwareManaged(r, 8);
```

```
        std::tie(pS, pSOffsets) = partition_softwareManaged(s, 8);
```

```
        break;
```

```
    }
```

```
// Step 2: partition-wise build & probe phase
```

```
    for (uint64_t i = 0, limit = pR.s.size(); i < limit; i++) {
        std::unordered_set<keyType> hashTable;
```

```
// Step 1: build phases
```

```
        auto start_r = pR.s[i];
```

```
        auto end_r = i == limit - 1 ? pR.r.size() : pR.s[i + 1];
```

```
        auto start_s = pS.s[i];
```

```
        auto end_s = i == limit - 1 ? pS.r.size() : pS.s[i + 1];
```

```
        if (p == Partitioning::softwareManaged) {
```

```
            end_r = std::min(end_r, start_r + pROffsets[i]);
```

```
            end_s = std::min(end_s, start_s + pSOffsets[i]);
```

```
        }
```

```
        for (uint64_t j = start_r; j < end_r; j++) {
```

```
            hashTable.insert(pR.r[j].key);
```

```
        }
```

```
// Step 2: probe phase
```

```
        for (uint64_t j = start_s; j < end_s; j++) {
```

```
            matches += hashTable.count(pS.r[j].key);
```

```
        }
```

```
    }
```

```
    return matches;
```

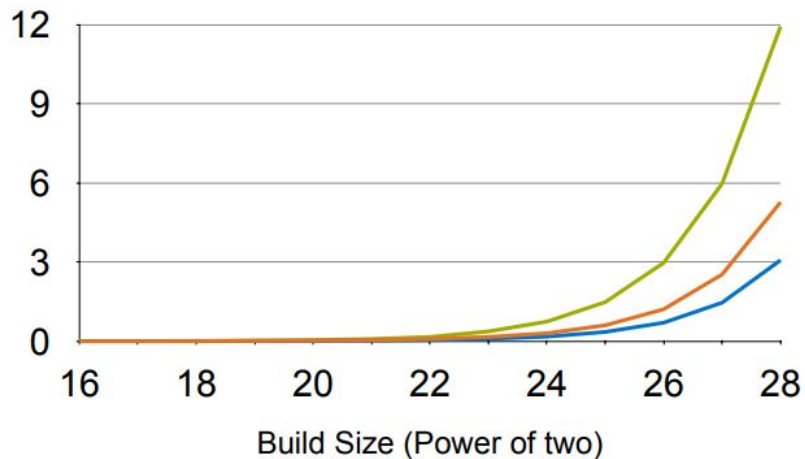
```
}
```

# Assignment 3


## Evaluation - Partitioning

- Partition Naive 8 bits
- Partition SMB 8 bits
- Partition Multipass 2x 2 bits (reused naive partitioning)

Time (s)



LLC-misses



TLB-misses



# Assignment 3 - Partitioning Evaluation

## Multi-Pass Partitioning (Green Line)

- Performs two full rounds of naive partitioning (coarse then fine), so it touches the entire dataset twice.
- As a result, it incurs double the histogram/build overhead and double the scatter traffic, causing the **longest overall runtime**.

**Software-Managed Buffers (SMB, Blue Line):** Buffers tuples in small, cache-line-sized chunks before writing them out. Keeps each 64-byte buffer hot in L1 while it fills. Uses non-temporal (streaming) stores to write full cache lines straight to memory, bypassing L1/L2.

- **Benefits:**
  - Far fewer cache-line evictions and write-allocate misses
  - Lower LLC and TLB miss rates (especially once data approaches LLC size)
- **Trade-Off:**
  - Slight branch-misprediction overhead for “is buffer full?” checks
  - More bookkeeping for per-partition counters and buffer flushes

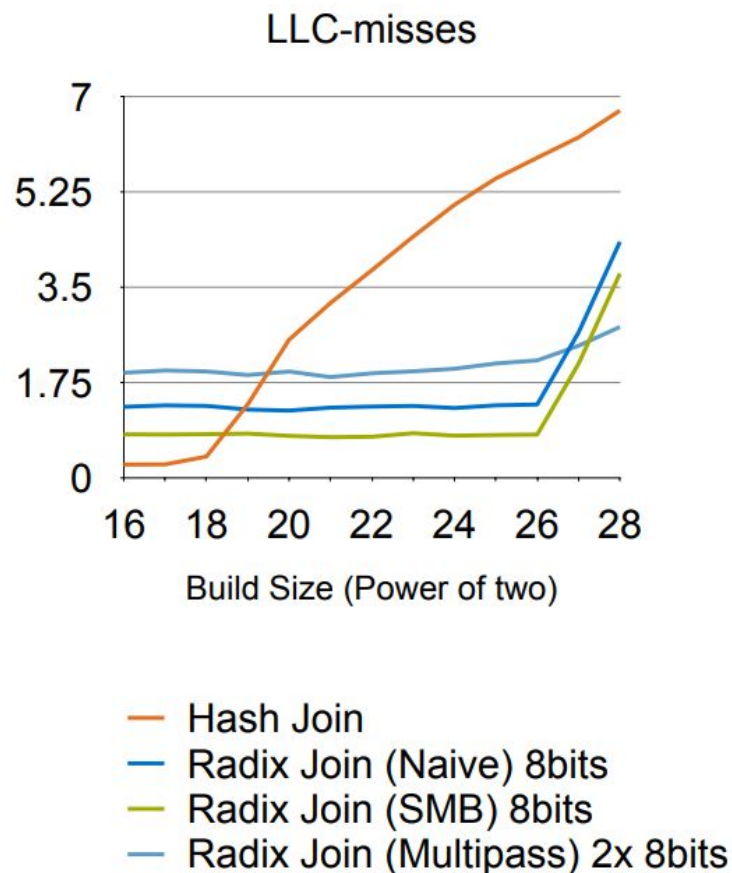
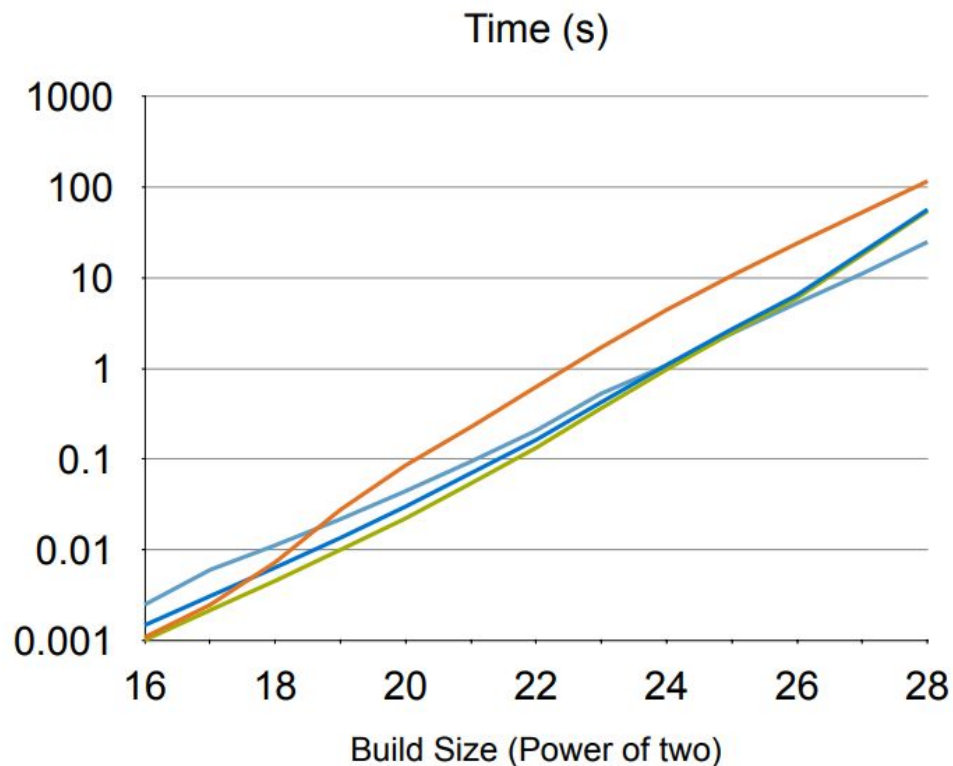
## Why SMB Beats Naïve Once Data Exceeds Cache

- Around  $2^{20} - 2^{21}$  (~1 – 2 million tuples  $\Rightarrow$  ~8 MB), working set nears our LLC (14 MB).
- **Naïve** scatter writes then forces repeated write-allocate and eviction of cached data, spiking LLC and TLB misses.
- **SMB** keeps write traffic in full 64-byte chunks, cutting LLC misses  $\rightarrow$  smoother, lower-latency writes.
- Consequently, SMB’s overall time curve remains below naïve once you exceed L2/L3 boundaries.

# Assignment 3

LLC-size: 14 MiB

## Evaluation - Join



# Assignment 3 - Radix Join Evaluation

**Hash Join: Fastest when the entire build side fits in cache (up to  $\sim 2^{18}$  tuples  $\approx 256$  KB)**

- No partitioning overhead—random hashtable lookups stay in L3.
- Once  $|R| > 2^{18}$ , LLC misses spike sharply as the hashtable no longer fits on-chip.

**Radix Join (Partition-then-Join): Partitions keep each bucket cache-resident until  $|R| \approx 2^{26}$  ( $\approx 64$  MB)**

- The “knee” in LLC misses appears around  $2^{18} + \text{\#bits}$  (i.e.  $2^{26}$  for 8-bit partitions).
- Below that, each bucket fits in L3/L2, so build+probe per bucket generates very few LLC misses.

## Partitioning Strategies

1. **Naïve (One-Pass):** Good once  $|R| > 2^{18}$ ; but as soon as partitions themselves exceed L3 ( $\approx 2^{26}$ ), LLC misses climb again.
2. **SMB (Software-Managed Buffers): Best in the mid-range ( $2^{18} \dots 2^{26}$ )**
  - Small, cache-line-sized write buffers reduce write-allocate traffic and keep LLC/TLB misses lowest.
  - After  $|R| \approx 2^{26}$ , partitions overflow L3 and LLC misses grow, but still below naïve one-pass.
3. **Multi-Pass (Two-Stage)**
  - **Only advantageous at the very largest sizes ( $\geq 2^{27}$ )**, because re-splitting partitions delays bucket overflow.
  - Before that point, the extra scan/overhead outweighs the benefit—LLC misses stay low but total time is higher than SMB.

# Questions?

