

# Data Processing on Modern Hardware

## Tutorial 8

Ferdinand Gruber

Michalis Georgoulakis

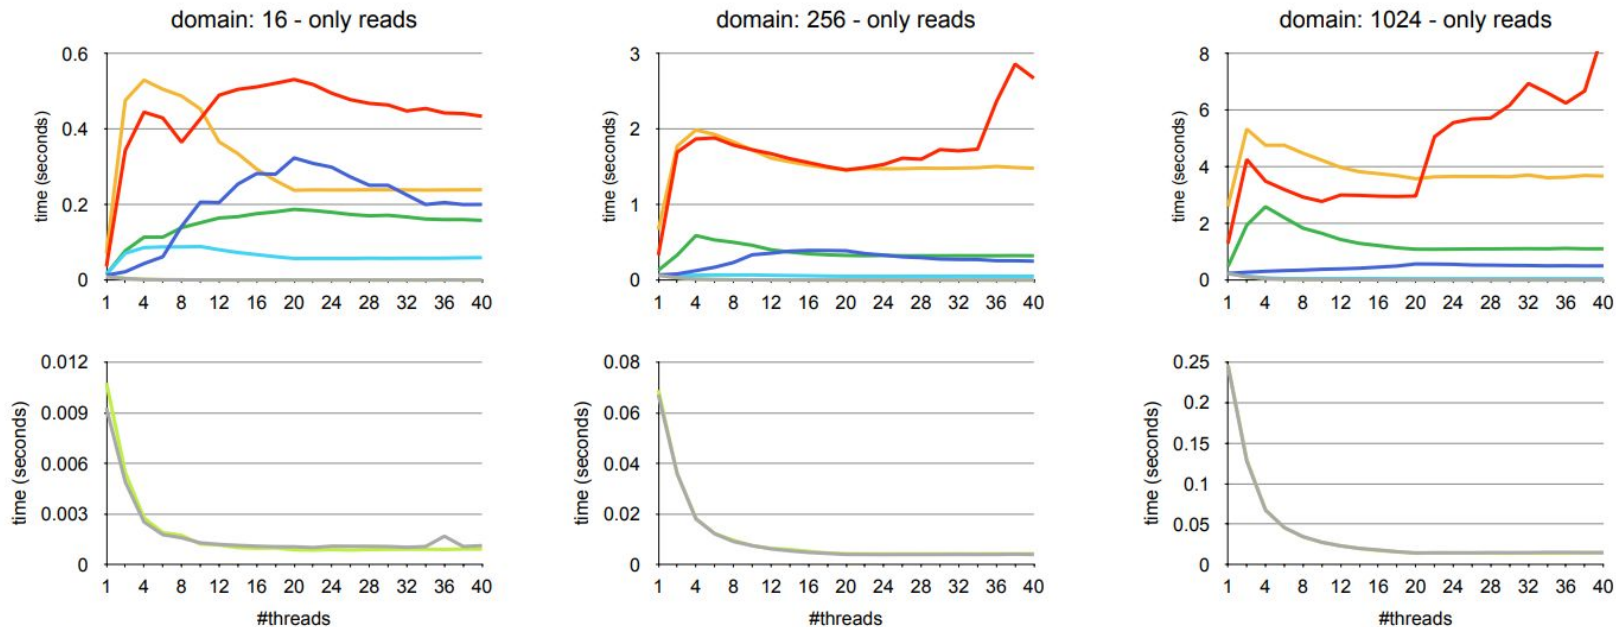


**We would be grateful if you take a few minutes to fill in the course evaluation you have received by email.**

# Assignment 5 - Synchronization







# Latency Comparison - Read-Only Workloads

- nosync
- coarse
- coarseRW
- lockCoupling
- lockCouplingRW
- optimistic
- optimisticLockCoupling









# Latency Comparison - Mixed Workloads

- nosync
- coarse
- coarseRW
- lockCoupling
- lockCouplingRW
- optimistic
- optimisticLockCoupling



# Perf Metrics

- nosync
- coarse
- coarseRW
- lockCoupling
- lockCouplingRW
- optimistic
- optimisticLockCoupling



# Conclusions

- Assignment Workload difficult to tune with synchronization as:
  - list operations are cheap
  - compared to the operations, locking is indeed expensive
- Fine-grained locking approaches produce a high overhead
  - especially lock coupling (unexpectedly), probably due to the number of locks - low IPC
  - coarse is decent for low domain: locking period is short
  - coarse-grained waiting times add up with multiple threads
- Optimistic Methods deliver the best performance among the synchronized methods
  - Optimistic Lock Coupling can even outperform single-threaded no sync approach
  - It also is the only method with quite low CPU cycles - minimal waiting overhead

# Optimistic Lock Coupling

*“The first mode is similar to a traditional mutex and excludes other threads by physically acquiring the underlying lock.”*

*“In the second mode, reads can proceed optimistically by validating a version counter that is embedded in the lock (similar to optimistic concurrency control). “*

The second mode can not be achieved by a conventional mutex. What we expected to see was the first mode, as exemplified in the previous tutorial's slides.



# Optimistic Lock Coupling: contains()

```

→ struct Entry {
    T key;
    Entry *next;
    uint64_t version = 0; // Version counter for the lock coupling
    M mutex;              // Mutex for each element in the list
};

Entry staticHead;
Entry staticTail;

```

```

bool contains(T k) {
    start:
    Entry *pred;
    Entry *curr = &staticHead;
→    uint64_t lastVersion = curr->version;

    while (curr->key < k) {
        pred = curr;                                // Make current the predecessor
        curr = curr->next;                            // Go to the next element
→        if (pred->version != lastVersion) // Validate using version
            goto start;
        lastVersion = curr->version; // Update for the new element
    }
    return (curr->key == k);
}

```

# Optimistic Lock Coupling: insert()

```
→ struct Entry {  
    T key;  
    Entry *next;  
    uint64_t version = 0; // Version counter for the lock coupling  
    M mutex;             // Mutex for each element in the list  
};  
  
Entry staticHead;  
Entry staticTail;
```

```
void insert(T k) {  
    start:  
    Entry *pred;  
    Entry *curr = &staticHead;  
→    uint64_t predVersion;  
    uint64_t currVersion = curr->version;  
  
    while (curr->key < k) {  
        pred = curr;                      // Make current the predecessor  
        curr = curr->next;                 // Go to the next element  
→    predVersion = currVersion;           // Make currVersion the predVersion  
        currVersion = curr->version;      // Update for the new element  
        if (pred->version != predVersion) // Validate using version  
            goto start;  
    }  
  
→    pred->mutex.lock(); // Lock the predecessor of the relevant element  
    curr->mutex.lock(); // Lock the relevant element  
  
→    if (pred->version != predVersion) {  
        pred->mutex.unlock();  
        curr->mutex.unlock();  
        goto start;  
    }  
  
    if (curr->key != k) {                  // Is the element in the list?  
        auto *n = new Entry{k, curr};    // Create the element  
        pred->next = n;                   // And insert it  
→    pred->version++;                      // Update the version due to changes  
    }  
→    pred->mutex.unlock();  
    curr->mutex.unlock();  
}
```

# Optimistic Lock Coupling

## remove()

```
→ struct Entry {  
    T key;  
    Entry *next;  
    uint64_t version = 0; // Version counter for the lock coupling  
    M mutex;             // Mutex for each element in the list  
};  
  
Entry staticHead;  
Entry staticTail;
```

```
void remove(T k) {  
    start:  
    Entry *pred;  
    Entry *curr = &staticHead;  
→    uint64_t predVersion;  
    uint64_t currVersion = curr->version;  
  
    while (curr->key < k) {  
        pred = curr;                                // Make current the predecessor  
        curr = curr->next;                            // Go to the next element  
→        predVersion = currVersion;                    // Make currVersion the predVersion  
        currVersion = curr->version;                  // Update for the new element  
        if (pred->version != predVersion) // Validate using version  
            goto start;  
    }  
  
→    pred->mutex.lock(); // Lock the predecessor of the relevant element  
    curr->mutex.lock(); // Lock the relevant element  
  
→    if (pred->version != predVersion) {  
        pred->mutex.unlock();  
        curr->mutex.unlock();  
        goto start;  
    }  
  
    if (curr->key == k) { // Is the element in the list?  
        pred->next = curr->next; // Then, remove it.  
→        pred->version++;        // Update the version due to changes  
        curr->version++;        // Update the version due to changes  
    }  
→    pred->mutex.unlock();  
    curr->mutex.unlock();  
}
```

# Landscape of Parallel Computing APIs



- **OpenCL**
  - Open standard by Khronos Group for heterogeneous computing
  - Portable across CPUs, GPUs, DSPs, FPGAs
- **CUDA**
  - NVIDIA-proprietary API for NVIDIA GPUs
  - Rich ecosystem (Thrust, cuDNN, etc.) and mature tooling
- **Vulkan Compute**
  - Explicit low-overhead graphics & compute API
  - Fine-grained control over memory and execution on GPUs
- **Other Approaches**
  - **SYCL**: Higher-level C++ wrapper over OpenCL
  - **DirectCompute**: Microsoft's compute API for Windows
- **When to Choose?**
  - Portability / Learning Curve: OpenCL/SYCL
  - Performance & Ecosystem: CUDA (on NVIDIA hardware)
  - Graphics + Compute Integration: Vulkan Compute

# OpenCL Execution Model

## Open Standard by Khronos Group

- Vendor-neutral API for parallel programming
- Widely adopted across hardware manufacturers (Gimp, Affinity, Sony Vegas, legacy Photoshop)

## Host vs. Device Code

- **Host:** C/C++ (or bindings) running on CPU
- **Device:** Kernels (C-like) compiled for GPU/accelerator

## Portability Across Architectures

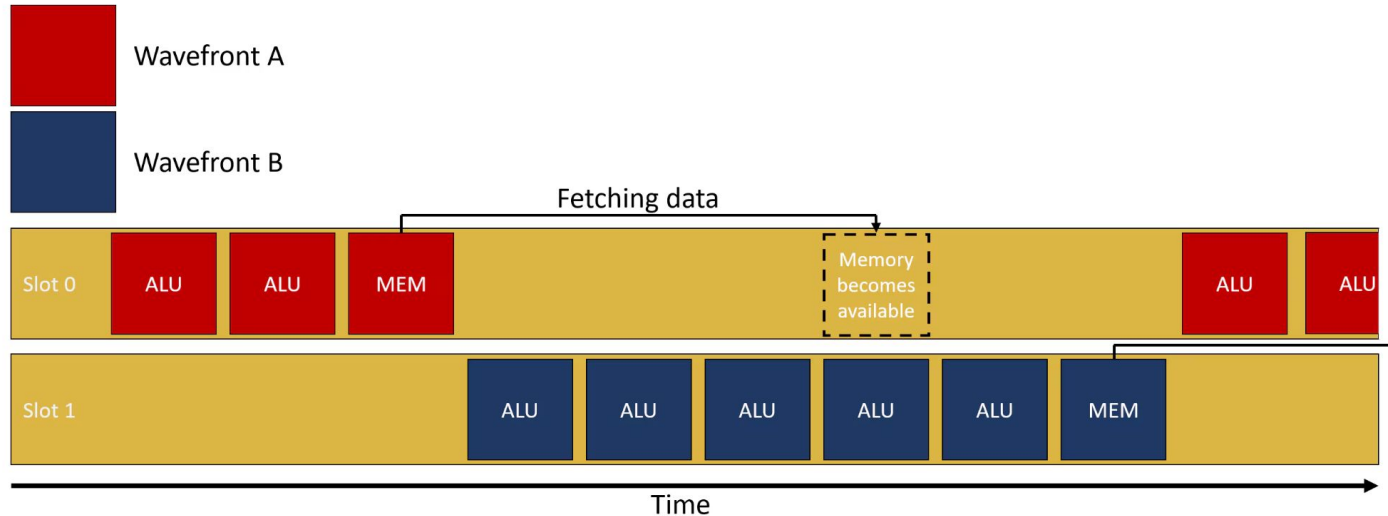
- Write once, run on CPUs, GPUs, DSPs, FPGAs
- Single source for heterogeneous platforms

# OpenCL Execution Model

- **Platform → Device → Context → Command Queue**
  - **Platform:** Implementation vendor (e.g., NVIDIA, AMD)
  - **Device:** Compute unit (GPU, CPU core)
  - **Context:** Manages objects (buffers, programs) across devices
  - **Command Queue:** Orders work for a specific device
- **Work-Items & Work-Groups**
  - **Work-Item:** Single execution instance of a kernel
    - i. Execute completely independently (aside from any explicit synchronization) and map directly to the hardware's SIMD lanes or scalar threads.
  - **Work-Group:** Group of work-items sharing local memory and synchronization
    - i. Execute on the same compute unit (CU) or multiprocessor
    - ii. Share a region of local memory (`__local`)
    - iii. Can synchronize with each other via `barrier(CLK_LOCAL_MEM_FENCE)`
- **NDRange**
  - Defines the grid ("global size") and grouping ("local size")

# Wavefront (AMD) / Warp (NVIDIA) Width

- GPUs don't really run each work-item entirely independently—they bundle them into fixed-size SIMD “lanes” that execute in lock-step:
- AMD calls this a **wavefront**, typically 64 work-items wide.
- NVIDIA calls it a **warp**, typically 32 threads wide.



# OpenCL Memory Model



## Global Memory

- Accessible by all work-items across all work-groups
- **Characteristics:** High latency (100s of cycles), large capacity (GBs)
- **Best Practice:**
  - Read large chunks coalesced (e.g., `float4` loads)
  - Minimize random accesses

## Constant Memory

- Read-only region of global memory, cached on device for fast lookups
- **Characteristics:** Low-latency on repeated reads, limited size (tens of KB)
- **Example Use:** Lookup tables (e.g., color conversion coefficients)

## Local Memory

- Shared among work-items in the **same** work-group
- **Characteristics:** Low latency (close to registers), limited size (tens of KB)

## Private Memory

- Per-work-item registers or stack space
- **Characteristics:** Fastest access, very limited size
- **Example Use:** Temporary accumulators or loop indices



## Data-Transfer Patterns

- **Host ↔ Global Memory** via explicit enqueues

```
clEnqueueWriteBuffer(queue, bufA, CL_TRUE, 0, size, hostA, 0, NULL, NULL);  
clEnqueueReadBuffer(queue, bufC, CL_TRUE, 0, size, hostC, 0, NULL, NULL);
```

- **Optimization Tips:**

- **Batch transfers** (fewer, larger copies rather than many small ones)
- **Asynchronous transfers** with events to overlap compute and copy
- **Use pinned (page-locked) host memory** for higher PCIe bandwidth

## Access Pattern Examples

- **Coalesced Access:** consecutive `get_global_id(0)` threads load contiguous floats → full bus utilization
- **Strided Access (Anti-pattern):** `data[i * stride]` with large `stride` → many small transactions, low throughput

# CUDA Syntax – Kernels and Memory

Defined by `__global__ <signature>`

```
// CUDA Kernel function to add the elements of two arrays on the GPU
__global__ void add(int n, float *x, float *y)
{
    int index = threadIdx.x;
    int stride = blockDim.x;
    for (int i = index; i < n; i += stride){
        y[i] = y[i] + x[i];
    }
}
```

Invoked with `<kernel_name><<<X,Y>>>`

```
// Run kernel on 1GB elements on the GPU
add<<<1, 256>>>>(N, gx, gy);
```

Memory allocation

- Static allocation `__device__ int X`
- Dynamic allocation `cudaMalloc(&X, size)`

```
// Initialize pointers and allocate memory
float *gx, *gy;
cudaMalloc((void*)&gx, N*sizeof(float));
cudaMalloc((void*)&gy, N*sizeof(float));
```

Memory copying (Host to Device, Device to Host)

- `cudaMemcpy`
- `cudaMemcpyManaged`

```
// Move data from CPU to GPU
cudaMemcpy(gx, hx, N*sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(gy, hy, N*sizeof(float), cudaMemcpyHostToDevice);

// Move data from GPU to CPU
cudaMemcpy(hy, gy, N*sizeof(float), cudaMemcpyDeviceToHost);
```

# Questions?

